

property_price_data_model

December 11, 2025

```
[ ]: # Problem Statement
# A major challenge for property sellers is determining the right sale price ↴
# for a property.
# Accurate price prediction benefits investors in assessing returns and helps ↴
# buyers plan
# their finances in line with market trends.
# Property prices are influenced by several factors such as:
# Property area
# Basement square footage
# Year of construction
# Number of bedrooms
# And many other features
# By applying Regression Analysis, we can model these features to predict the ↴
# price of
# a property with better accuracy.
```

```
[27]: # import libraries:
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import joblib
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
[28]: # load the data
data = pd.read_csv("property_price_data (1).csv")
data.head(10)
```

```
[28]:      \ MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape \
0    PR0504        20       RL     100.0   15537    Pave    NaN    IR1
1    PR0102        60       RL      77.0    9534    Pave    NaN    Reg
2    PR0609        70       RL      NaN    12781    Pave    NaN    Reg
3    PR01090       120      FV      37.0    3728    Pave    Pave   IR1
4    PR0820       120       RL      44.0    6606    Pave    NaN    IR1
```

```

5   PR0685          60      RL      58.0    18002    Pave     NaN     IR2
6   PR01281         20      RL      67.0     9769    Pave     NaN     IR1
7   PR0921          60      RL      70.0     8653    Pave     NaN     IR1
8   PR01454         20      RL      90.0    17720    Pave     NaN     Reg
9   PR0541          20      RL      85.0    15369    Pave     NaN     Reg

```

```

LandContour Utilities ... 3SsnPorch ScreenPorch PoolArea PoolQC Fence \
0      Lvl    AllPub ...        0       161        0     NaN  GdWo
1      Lvl    AllPub ...        0        0        0     NaN  NaN
2      HLS    AllPub ...        0        0        0     NaN  NaN
3      Lvl    AllPub ...        0        0        0     NaN  NaN
4      Lvl    AllPub ...        0        0        0     NaN  NaN
5      Lvl    AllPub ...        0        0        0     NaN  NaN
6      Lvl    AllPub ...        0        0        0     NaN  NaN
7      Lvl    AllPub ...        0        0        0     NaN  NaN
8      Lvl    AllPub ...        0        0        0     NaN  NaN
9      Lvl    AllPub ...        0        0        0     NaN  NaN

```

```

MiscFeature MiscVal YrSold SaleCondition SalePrice
0      NaN      0    2010      Normal    288330
1      NaN      0    2010      Normal    183164
2      NaN      0    2007    Alloca    362145
3      NaN      0    2006      Normal    196079
4      NaN      0    2010    Partial    228515
5      NaN      0    2010      Normal    224119
6      NaN      0    2009      Normal    228107
7      NaN      0    2007      Normal    203953
8      NaN      0    2006    Abnorml    87468
9      NaN      0    2009      Normal    313697

```

[10 rows x 69 columns]

[3]: data.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 970 entries, 0 to 969
Data columns (total 69 columns):
 #   Column           Non-Null Count  Dtype  
 ____ _-----_ 
 0   \              970 non-null    object 
 1   MSSubClass       970 non-null    int64  
 2   MSZoning        970 non-null    object 
 3   LotFrontage     789 non-null    float64
 4   LotArea         970 non-null    int64  
 5   Street          970 non-null    object 
 6   Alley            56 non-null    object 
 7   LotShape        970 non-null    object 
 8   LandContour     970 non-null    object 

```

9	Utilities	970	non-null	object
10	LotConfig	970	non-null	object
11	LandSlope	970	non-null	object
12	Neighborhood	970	non-null	object
13	Condition1	970	non-null	object
14	Condition2	970	non-null	object
15	BldgType	970	non-null	object
16	PropStyle	970	non-null	object
17	OverallQual	970	non-null	int64
18	OverallCond	970	non-null	int64
19	YearBuilt	970	non-null	int64
20	YearRemodAdd	970	non-null	int64
21	RoofStyle	970	non-null	object
22	RoofMatl	970	non-null	object
23	Exterior1st	970	non-null	object
24	Exterior2nd	970	non-null	object
25	ExterQual	970	non-null	object
26	ExterCond	970	non-null	object
27	Foundation	970	non-null	object
28	BsmtQual	947	non-null	object
29	BsmtCond	947	non-null	object
30	BsmtExposure	946	non-null	object
31	TotalBsmtSF	970	non-null	int64
32	Heating	970	non-null	object
33	HeatingQC	970	non-null	object
34	CentralAir	970	non-null	object
35	Electrical	970	non-null	object
36	GrLivArea	970	non-null	int64
37	BsmtFullBath	970	non-null	int64
38	BsmtHalfBath	970	non-null	int64
39	FullBath	970	non-null	int64
40	HalfBath	970	non-null	int64
41	BedroomAbvGr	970	non-null	int64
42	KitchenAbvGr	970	non-null	int64
43	KitchenQual	970	non-null	object
44	TotRmsAbvGrd	970	non-null	int64
45	Functional	970	non-null	object
46	Fireplaces	970	non-null	int64
47	FireplaceQu	533	non-null	object
48	GarageType	918	non-null	object
49	GarageYrBlt	918	non-null	float64
50	GarageFinish	918	non-null	object
51	GarageCars	970	non-null	int64
52	GarageArea	970	non-null	int64
53	GarageQual	918	non-null	object
54	GarageCond	918	non-null	object
55	PavedDrive	970	non-null	object
56	WoodDeckSF	970	non-null	int64

```
57 OpenPorchSF    970 non-null    int64
58 EnclosedPorch  970 non-null    int64
59 3SsnPorch      970 non-null    int64
60 ScreenPorch    970 non-null    int64
61 PoolArea       970 non-null    int64
62 PoolQC         7 non-null     object
63 Fence          169 non-null    object
64 MiscFeature    29 non-null     object
65 MiscVal        970 non-null    int64
66 YrSold         970 non-null    int64
67 SaleCondition   970 non-null    object
68 SalePrice      970 non-null    int64
dtypes: float64(2), int64(27), object(40)
memory usage: 523.0+ KB
```

```
[4]: data.shape
```

```
[4]: (970, 69)
```

```
[5]: data.isnull().sum()
```

```
[5]: \
MSSubClass          0
MSZoning            0
LotFrontage         181
LotArea              0
...
MiscFeature         941
MiscVal              0
YrSold                0
SaleCondition         0
SalePrice              0
Length: 69, dtype: int64
```

```
[6]: data[data.isnull().any(axis=1)]      # show the rows that is having null values
```

```
[6]:      \ MSSubClass MSZoning  LotFrontage  LotArea Street Alley LotShape \
0      PR0504      20      RL      100.0    15537  Pave  NaN  IR1
1      PR0102      60      RL      77.0     9534  Pave  NaN  Reg
2      PR0609      70      RL      NaN     12781  Pave  NaN  Reg
3      PR01090     120     FV      37.0     3728  Pave  Pave  IR1
4      PR0820      120     RL      44.0     6606  Pave  NaN  IR1
...
965     PR0162      60      RL      110.0    14777  Pave  NaN  IR1
966     PR0693      60      RL      42.0     26949  Pave  NaN  IR1
967     PR0674      20      RL      110.0    15573  Pave  NaN  Reg
968     PR0750      50      RL      50.0     8367  Pave  NaN  Reg
969     PR01032     75      RL      102.0    16765  Pave  NaN  Reg
```

```

LandContour Utilities ... 3SsnPorch ScreenPorch PoolArea PoolQC Fence \
0      Lvl    AllPub ...     0      161      0    NaN  GdWo
1      Lvl    AllPub ...     0      0      0    NaN  NaN
2      HLS    AllPub ...     0      0      0    NaN  NaN
3      Lvl    AllPub ...     0      0      0    NaN  NaN
4      Lvl    AllPub ...     0      0      0    NaN  NaN
...
965     Lvl    AllPub ...     0      0      0    NaN  NaN
966     Lvl    AllPub ...     0      0      0    NaN  NaN
967     Lvl    AllPub ...     0      200     0    NaN  NaN
968     Lvl    AllPub ...     0      0      0    NaN  NaN
969     Lvl    AllPub ...     0      0      0    NaN  NaN

MiscFeature MiscVal YrSold SaleCondition SalePrice
0      NaN      0    2010      Normal   288330
1      NaN      0    2010      Normal   183164
2      NaN      0    2007      Allocat  362145
3      NaN      0    2006      Normal   196079
4      NaN      0    2010      Partial  228515
...
965     NaN      0    2008      Normal   413403
966     NaN      0    2006      Normal   336805
967     NaN      0    2007      Normal   259324
968     NaN      0    2009      Normal   96478
969     NaN      0    2009      Normal   197073

```

[970 rows x 69 columns]

```
[7]: data.duplicated().sum()
```

```
[7]: np.int64(0)
```

```
[29]: # filling the missing values
```

```
# fill numeric columns with median
data = data.fillna(data.median(numeric_only=True))

# fill categorical columns with mode
for col in data.select_dtypes(include='object'):
    data[col] = data[col].fillna(data[col].mode()[0])
```

```
[30]: data.isnull().sum()
```


```
[30]: \
MSSubClass      0
MSZoning       0
```

```
LotFrontage      0
LotArea          0
...
MiscFeature      0
MiscVal          0
YrSold           0
SaleCondition    0
SalePrice         0
Length: 69, dtype: int64
```


```
[69]: # outlier detection and handling:
```

```
numeric_cols = data.select_dtypes(include=['int64', 'float64']).columns


for col in numeric_cols:
    plt.figure(figsize=(6,3))
    sns.boxplot(x=data[col])
    plt.title(f"Boxplot of {col}")
    plt.show()
```




Boxplot of LotFrontage




Boxplot of LotArea




Boxplot of OverallQual




Boxplot of OverallCond




Boxplot of YearBuilt




Boxplot of YearRemodAdd




Boxplot of TotalBsmtSF




Boxplot of GrLivArea




Boxplot of BsmtFullBath




Boxplot of BsmtHalfBath




Boxplot of FullBath




Boxplot of HalfBath




Boxplot of BedroomAbvGr




Boxplot of KitchenAbvGr




Boxplot of TotRmsAbvGrd




Boxplot of Fireplaces




Boxplot of GarageYrBlt




Boxplot of GarageCars




Boxplot of GarageArea




Boxplot of WoodDeckSF




Boxplot of OpenPorchSF




Boxplot of EnclosedPorch




Boxplot of 3SsnPorch




Boxplot of ScreenPorch





Boxplot of PoolArea



Boxplot of MiscVal





```
[70]: # removing outliers
```

```
def remove_outliers_iqr(data, cols):
    cleaned = data.copy()
    for col in cols:
        Q1 = cleaned[col].quantile(0.25)
        Q3 = cleaned[col].quantile(0.75)
```

```

IQR = Q3 - Q1

lower = Q1 - 1.5*IQR
upper = Q3 + 1.5*IQR

cleaned = cleaned[(cleaned[col] >= lower) & (cleaned[col] <= upper)]
return cleaned

df_clean = remove_outliers_iqr(data, numeric_cols)
print(df_clean.shape)

```

(421, 1180)

[54]: # converting categorical columns into numeric

```

data_model = pd.get_dummies(df_clean, drop_first=True)

data_model.head()

```

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	\
1	60	77.0	9534	6	5	1985	
3	120	37.0	3728	8	5	2005	
4	120	44.0	6606	7	5	2009	
5	60	58.0	18002	7	5	1998	
6	20	67.0	9769	7	5	2002	
	YearRemodAdd	TotalBsmtSF	GrLivArea	BsmtFullBath	...	Fence_MnPrv	\
1	1985	741	1732	0	...	True	
3	2005	1247	1247	1	...	True	
4	2010	1358	1358	1	...	True	
5	1998	1195	1839	0	...	True	
6	2002	1573	1573	1	...	True	
	Fence_MnWw	MiscFeature_Othr	MiscFeature_Shed	MiscFeature_TenC	\		
1	False	False	True	False			
3	False	False	True	False			
4	False	False	True	False			
5	False	False	True	False			
6	False	False	True	False			
	SaleCondition_AdjLand	SaleCondition_Alloca	SaleCondition_Family	\			
1	False	False	False				
3	False	False	False				
4	False	False	False				
5	False	False	False				
6	False	False	False				
	SaleCondition_Normal	SaleCondition_Partial					

```

1          True      False
3          True      False
4         False     True
5          True      False
6          True      False

```

[5 rows x 1180 columns]

```

[55]: # feature correlation analysis: understand which numeric features strongly
      ↳correlate with
      # the target SalePrice.

corr = data_model.corr()

# sort correlation of all features with Saleprice
corr_with_target = corr["SalePrice"].sort_values(ascending=False)
print(corr_with_target.head(10))

```

```

SalePrice      1.000000
OverallQual    0.841627
GrLivArea      0.781242
GarageCars      0.705403
GarageArea      0.685849
YearBuilt       0.663157
FullBath        0.645180
GarageYrBlt     0.624319
Foundation_PConc 0.617840
TotRmsAbvGrd    0.607961
Name: SalePrice, dtype: float64

```

```

[44]: # Explaination of feature correlation analysis:

# higher correlation - stronger impact on price
# feature like overallqual, grlivearea, garagecars usually rank highest

```

```

[68]: # define independent(x) and dependent(y) variable

x = data_model.drop("SalePrice", axis=1)
y = data_model["SalePrice"]

```

```

[66]: # split the data

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
      ↳random_state=42)

# initialise the linear regression model on the training dataset

model = LinearRegression()

```

```
model.fit(x_train, y_train)

print("model training completed")
```

```
model training completed
```

```
[67]: # predicting on the test data

y_pred = model.predict(x_test)

# evaluate the model
mae = mean_absolute_error(y_test,y_pred)
mse = mean_squared_error(y_test,y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test,y_pred)

# performance metrics calculation
print("model evaluation metrics:")
print(f"mean absolute error (mae): {mae:.2f}")
print(f"mean squared error (mse): {mse:.2f}")
print(f"root measured squared error (rmse): {rmse:.2f}")
print(f"r2 score: {r2:.2f}\n")
```

```
model evaluation metrics:
mean absolute error (mae): 15593.09
mean squared error (mse): 453635879.40
root measured squared error (rmse): 21298.73
r2 score: 0.87
```

```
[59]: # check important features (coefficients)

importance = pd.DataFrame({
    'Feature' : x.columns,
    'Coefficient': model.coef_
}).sort_values(by='Coefficient', ascending=False)

importance.head(10)
```

```
[59]:      Feature   Coefficient
320     \_PR014  55229.652942
921     \_PR0886  33039.479336
769     \_PR0681  28839.074153
913     \_PR0878  28051.299698
75      \_PR0108  27700.936378
174     \_PR01213 26704.940316
772     \_PR0684  25944.913291
42      \_PR01028 24514.730709
```

```
264 \_PR01320 24279.832532
68 \_PR01066 24256.777782
```

[60]: # graphical evaluation for model performance

```
# predicted vs actual plot
plt.figure(figsize=(8,6))
plt.scatter(y_test, y_pred, alpha=0.8, edgecolor="k")

line = np.linspace(min(y_test.min(), y_pred.min()), max(y_test.max(), y_pred.
    ↴max()), 100)
plt.plot(line, line, linewidth=2)

plt.xlabel("actual price (lakhs)")
plt.ylabel("predicted price (lakhs)")
plt.grid(True)
plt.title("predicted vs actual property prices")
plt.show()
```



[16]: # Business interpretation:

- # 1. Strong Positive Relationship (Model is Learning Well)
 - # Most points are clustered close to the diagonal line, which represents ↴perfect predictions.
 - # This indicates that the model is able to capture the overall pricing pattern ↴correctly.
- # 2. Good Performance for Mid-Range Properties
 - # For properties priced between 1 lakh to 3.5 lakh (100k - 350k), the ↴predicted values closely match the actual values, showing:
 - # Consistent model behavior
 - # Low prediction error
 - # Reliable accuracy in the common price range
- # 3. Slight Under/Over-Prediction at Higher Prices
 - # For higher-priced properties (above 4 lakh / 400k), the scatter points begin ↴moving away from the ideal line, which means:
 - # The model may under-predict some expensive properties
 - # A few points are far from the line → higher error
 - # The linear model may not fully capture the non-linear behavior at high ↴price ranges
 - # This is normal when using a simple linear regression model.
- # 4. Business Insight
 - # Overall, the model can be used for:
 - # Property price estimation
 - # Basic valuation
 - # Pricing decisions for average-priced properties
 - # Market trend analysis

[]: # Business Recommendations

- # Based on the model insights and feature relationships, here are practical ↴business recommendations:
 - # Focus on the Top Price-Driving Features
 - # The model shows that certain property features have the strongest correlation ↴with price (e.g., area, bedrooms, location score, property age, amenities).
 - # Recommendation:
 - # Real-estate companies should prioritize acquiring and marketing properties ↴that score high on these features since they significantly increase selling ↴price.
 - # Improve Listings With High-Impact Features
 - # Properties with certain high-correlation features can be made more valuable.

```

# Examples:
# Add amenities (parking, lift, security)
# Renovate older properties
# Improve interior quality
# These upgrades can increase predicted selling price, leading to better ↴ returns.

# Identify Undervalued Properties
# Since the model predicts a fair price:
# Properties selling below the model prediction can be targeted as "undervalued" ↴ opportunities".
# Investors can purchase and renovate these for profit.

# Target Marketing Based on Customer Budget
# The model helps predict price ranges.
# Real-estate companies can:
# Match customers to properties that fit their affordability
# Reduce search time
# Improve conversion rates

# Use Model Predictions for Better Negotiation
# The model provides a data-driven benchmark price.
# Agents can:
# Use predicted price as a reference in negotiations
# Avoid overpricing and losing customers
# Avoid underpricing and losing revenue

```

[]: # Model Limitations

```

# Linear Regression only captures linear relationships
# Real-estate pricing often depends on non-linear factors (e.g., large area ↴ increments increase price exponentially).
# Linear regression oversimplifies these relationships + reducing accuracy.

# Sensitive to Outliers
# Your model performance changed significantly before/after outlier removal.
# Linear Regression is not robust + extreme values distort coefficients.

# Multicollinearity may still exist
# Some features may be highly correlated (e.g., area, rooms).
# This makes model coefficients unstable and affects interpretability.

# Dataset may not include all price-driving factors
# Missing real-world factors:
# Distance to metro
# Crime rate
# Builder reputation

```

```

# Floor number
# Road width
# Because these are absent, the model cannot fully capture true price behavior.

# Scaling did not improve performance, after scaling the value of r2 is ↴
# decreasing to 0.82.
# This confirms that feature distributions are not ideal, and the data may be ↴
# skewed.

# Limited generalization
# If the dataset covers only a specific region (city/locality), the model ↴
# cannot be applied to other regions without retraining.

```

[72]: # Future Work / Improvements

```

# Try Advanced Models (Non-linear & More Accurate)
# To achieve better accuracy, test:
# Random Forest Regressor
# XGBoost
# Gradient Boosting Regressor
# These models handle non-linearity and outliers far better than Linear Regression.
# Accuracy can increase to 0.93-0.96.

# Build Feature Engineering Pipeline
# Improve features by adding:
# Price per square foot
# Age category (0-5 yrs, 5-15 yrs, 15+ yrs)
# Location rating
# Number of amenities
# Furnishing level category
# Better features → better predictions.

# Use Regularization Techniques
# To reduce multicollinearity:
# Ridge Regression
# Lasso Regression
# ElasticNet
# These make the model more stable and interpretable.

# Add More Real-Estate Data
# Add important features:
# Distance to schools
# Locality tier (A/B/C)
# Crime/safety score
# Connectivity score
# Floor number

```

```
# Facing direction  
# These can significantly improve R2.  
  
# Build a Residual Diagnostics Report  
# Check:  
# Homoscedasticity  
# Normality of residuals  
# Autocorrelation  
# These ensure the model meets all assumptions.
```

[]: